Can reasoning models bound an unlisted holding?

First responses and a separate technical recovery series

Gemini 3.1 Pro at High thinking returned 14 of 20 original answers. All 14 gave the supported answer type and exactly matched the reference numbers. One technical replay of each of the six missing items recovered four more answers. Two interval items still lack a visible answer.

Correct answers among observed responses

	Point controls	Tight ranges	Missing / 20
Qwen3-8B thinking, first responses	10/10	0/10	0
GPT-6 High, first responses	10/10	10/10	0
Gemini 3.1 Pro High, first responses	8/8	6/6	6
Gemini, including technical recovery	10/10	8/8	2

Ten point and ten interval questions were scheduled per model. The last row includes the first-response row. Primary tolerance was 0.005 percentage points per value and strict tolerance was 0.001. Both give the displayed counts. All 18 observed Gemini account answers are numerically exact. On interval items Qwen asserted a point in all ten replies; none of the eight recovered Gemini replies did so.

What happened during collection

The signed-in account exposed the registered model and High setting. In separate unscored checks, Gemini Pro accepted a one-word request and returned to its welcome screen without an answer. Flash returned a generic provider error on the same request. Subsequent Pro calls produced full answers. This establishes intermittent delivery in the test browser. The underlying technical cause remains unknown. P01, P02, P05 and P08 returned answers only on replay. P03 and P04 yielded none on either attempt.

What the experiment supports

On P09, both hosted models gave the attainable account interval of 1.83% to 2.81%. Qwen asserted 1.83%. The withheld Vanguard row implies 2.46% for the constructed account. Scoring tests the answer type and endpoints. It does not certify every sentence of reasoning. For example, Gemini described ranks as strictly descending on P15 although the prompt permitted ties; it still supplied the correct endpoints.

Decision

The prespecified stopping rule requires at least nine correct first answers in each group for each hosted model. Six original Gemini outcomes are missing, including four interval questions. The original stopping rule therefore remains unresolved. The observed answers give no positive signal for expanding the clean task into a general claim of strong-model failure. Generalisation beyond these configurations, ten constructed accounts and the stated ranking assumptions requires new evidence.